



NEW YORK UNIVERSITY

Backpropagation

<http://bit.ly/DLSP20>

Yann LeCun

NYU - Courant Institute & Center for Data Science

Facebook AI Research

<http://yann.lecun.com>

TAs: Alfredo Canziani, Mark Goldstein

Deep Learning, NYU, Spring 2020

Parameterized Model

► Parameterized model

$$\bar{y} = G(x, w)$$


► Example: linear regression

$$\bar{y} = \sum_i w_i x_i \quad C(y, \bar{y}) = ||y - \bar{y}||^2$$

► Example: Nearest neighbor:

$$\bar{y} = \operatorname{argmin}_k ||x - w_{k, \cdot}||^2$$

► Computing function G may involve complicated algorithms



Block diagram notations for computation graphs

▶ Variables (tensor, scalar, continuous, discrete...)




▶ Observed: input, desired output...



▶ Computed variable: outputs of deterministic functions


▶ Deterministic function



▶ Multiple inputs and outputs (tensors, scalars,...)

▶ Implicit parameter variable (here: w)

▶ Scalar-valued function (implicit output)



▶ Single scalar output (implicit)

▶ used mostly for cost functions

Loss function, average loss.

► Simple per-sample loss function


$$L(x, y, w) = C(y, G(x, w))$$

► A set of samples

$$S = \{(x[p], y[p]) \mid p = 0 \dots P - 1\}$$

► Average loss over the set

$$\mathcal{L}(S, w) = \frac{1}{P} \sum_{(x,y)} L(x, y, w) = \frac{1}{P} \sum_{p=0}^{P-1} L(x[p], y[p], w)$$



Gradient Descent

► Full (batch) gradient

$$w \leftarrow w - \eta \frac{\partial \mathcal{L}(S, w)}{\partial w}$$






► Stochastic Gradient (SGD)

- Pick a p in $0 \dots P-1$, then update w :

$$w \leftarrow w - \eta \frac{\partial L(x[p], y[p], w)}{\partial w}$$


► SGD exploits the redundancy in the samples

- It goes faster than full gradient in most cases
- In practice, we use mini-batches for parallelization.



Traditional Neural Net


- ▶ **Stacked linear and non-linear functional blocks**
 - ▶ Weighted sums, matrix-vector product
 - ▶ Point-wise non-linearities (e.g. ReLu, tanh,)



Traditional Neural Net

► Stacked linear and non-linear functional blocks

$$s[i] = \sum_{j \in \text{UP}(i)} w[i, j] \cdot z[j] \quad z[i] = f(s[i])$$



Backprop through a non-linear function

► Chain rule:

$$g(h(s))' = g'(h(s)).h'(s)$$

$$dc/ds = dc/dz * dz/ds$$

$$dc/ds = dc/dz * h'(s)$$

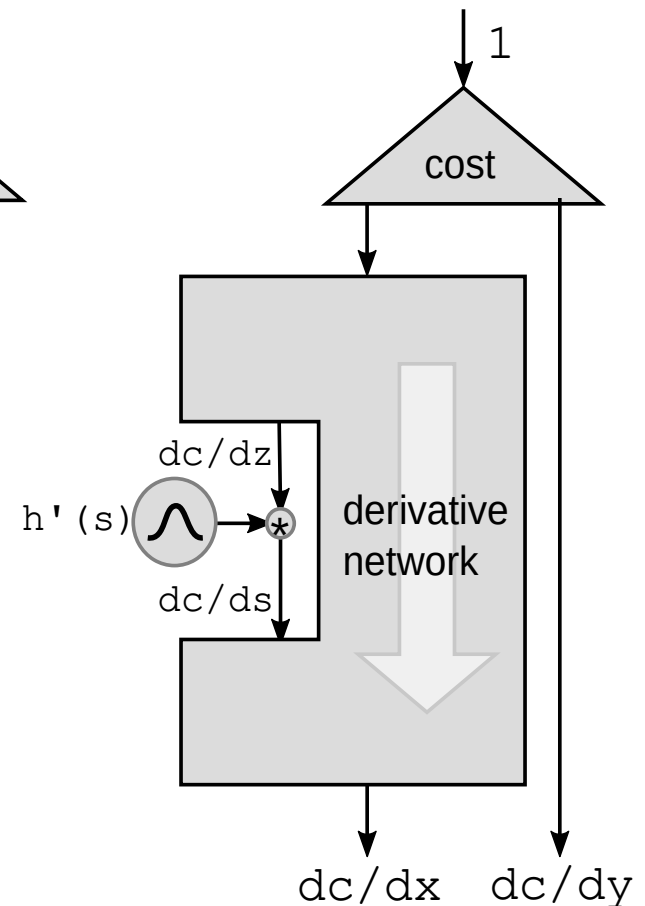

► Perturbations:

- Perturbing s by ds will perturb z by: $dz = ds * h'(s)$

- This will perturb c by

$$dc = dz * dc/dz = ds * h'(s) * dc/dz$$

- Hence: $dc/ds = dc/dz * h'(s)$





Backprop through a weighted sum

► Perturbations:

- Perturbing z by dz will perturb $s[0], s[1], s[2]$ by $ds[0]=w[0]*dz$, $ds[1]=w[1]*dz$, $ds[2]=w[2]*dz$
- This will perturb c by


$$dc = ds[0]*dc/ds[0] + ds[1]*dc/ds[1] + ds[2]*dc/ds[2]$$
- Hence: $dc/dz = dc/ds[0]*w[0] + dc/ds[1]*w[1] + dc/ds[2]*w[2] +$



Block Diagram of a Traditional Neural Net

▶ linear blocks $s_{k+1} = w_k z_k$

▶ Non-linear blocks $z_k = h(s_k)$



PyTorch definition

► Object-oriented version

- Uses predefined nn.Linear class, (which includes a bias vector)
- Uses torch.relu function
- State variables are temporary

```
import torch
```

```
from torch import nn
```

```
image = torch.randn(3, 10, 20)
```

```
d0 = image.nelement()
```

```
class mynet(nn.Module):
```

```
    def __init__(self, d0,d1,d2,d3):
```

```
        super().__init__()
```

```
        self.m0 = nn.Linear(d0, d1)
```

```
        self.m1 = nn.Linear(d1, d2)
```

```
        self.m2 = nn.Linear(d2, d3)
```

```
    def forward(self, x):
```

```
        z0 = x.view(-1)  ## flatten input tensor
```

```
        s1 = self.m0(x)
```

```
        z1 = torch.relu(s1)
```

```
        s2 = self.m1(z1)
```


```
        z2 = torch.relu(s2)
```

```
        s3 = self.m2(z2)
```

```
        return s3
```

```
model = mynet(d0,60,40,10)
```

```
out = model(image)
```



Backprop through a functional module

► Using chain rule for vector functions

$$z_g : [d_g \times 1] \quad z_f : [d_f \times 1]$$


$$\frac{\partial c}{\partial z_f} = \frac{\partial c}{\partial z_g} \frac{\partial z_g}{\partial z_f}$$

$$[1 \times d_f] = [1 \times d_g] * [d_g \times d_f]$$

► Jacobian matrix

► Partial derivative of i-th output w.r.t. j-th input

$$\left(\frac{\partial z_g}{\partial z_f} \right)_{ij} = \frac{(\partial z_g)_i}{(\partial z_f)_j}$$



Backprop through a multi-stage graph


► Using chain rule for vector functions

$$\frac{\partial c}{\partial z_k} = \frac{\partial c}{\partial z_{k+1}} \frac{\partial z_{k+1}}{\partial z_k} = \frac{\partial c}{\partial z_{k+1}} \frac{\partial f_k(z_k, w_k)}{\partial z_k}$$

$$\frac{\partial c}{\partial w_k} = \frac{\partial c}{\partial z_{k+1}} \frac{\partial z_{k+1}}{\partial w_k} = \frac{\partial c}{\partial z_{k+1}} \frac{\partial f_k(z_k, w_k)}{\partial w_k}$$

► Two Jacobian matrices for the module:

- One with respect to $z[k]$
- One with respect to $w[k]$



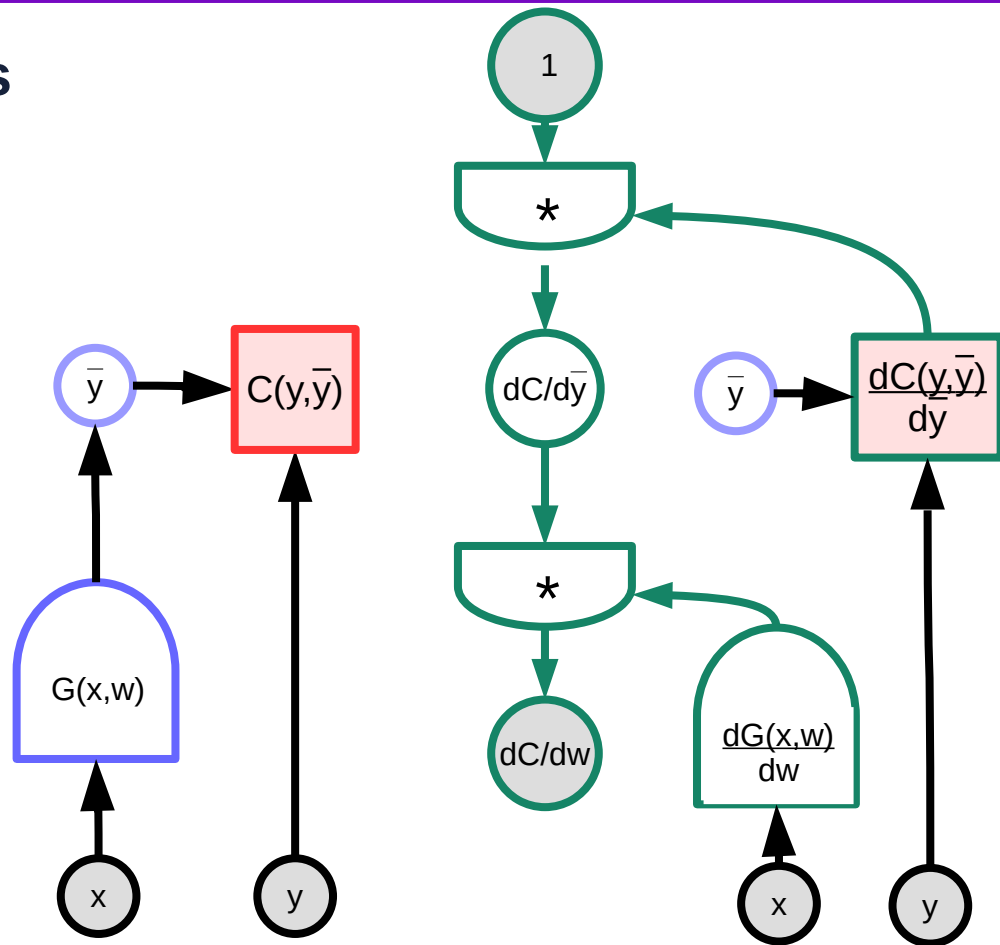
Backprop = propagation through a transformed graph

► Derivative of composed functions

$$C(G(w))' = C'(G(w))G'(w)$$

$$\frac{\partial C(y, \bar{y})}{\partial w} = \frac{\partial C(y, \bar{y})}{\partial \bar{y}} \frac{\partial \bar{y}}{\partial w}$$

$$\frac{\partial C(y, \bar{y})}{\partial w} = \frac{\partial C(y, \bar{y})}{\partial \bar{y}} \frac{\partial G(x, w)}{\partial w}$$



Gradient, Jacobian,

► Dimensions:

$$y, \bar{y} : [M \times 1] \quad w : [N \times 1]$$

$$\frac{\partial C(y, \bar{y})}{\partial w} = \frac{\partial C(y, \bar{y})}{\partial \bar{y}} \frac{\partial \bar{y}}{\partial w}$$

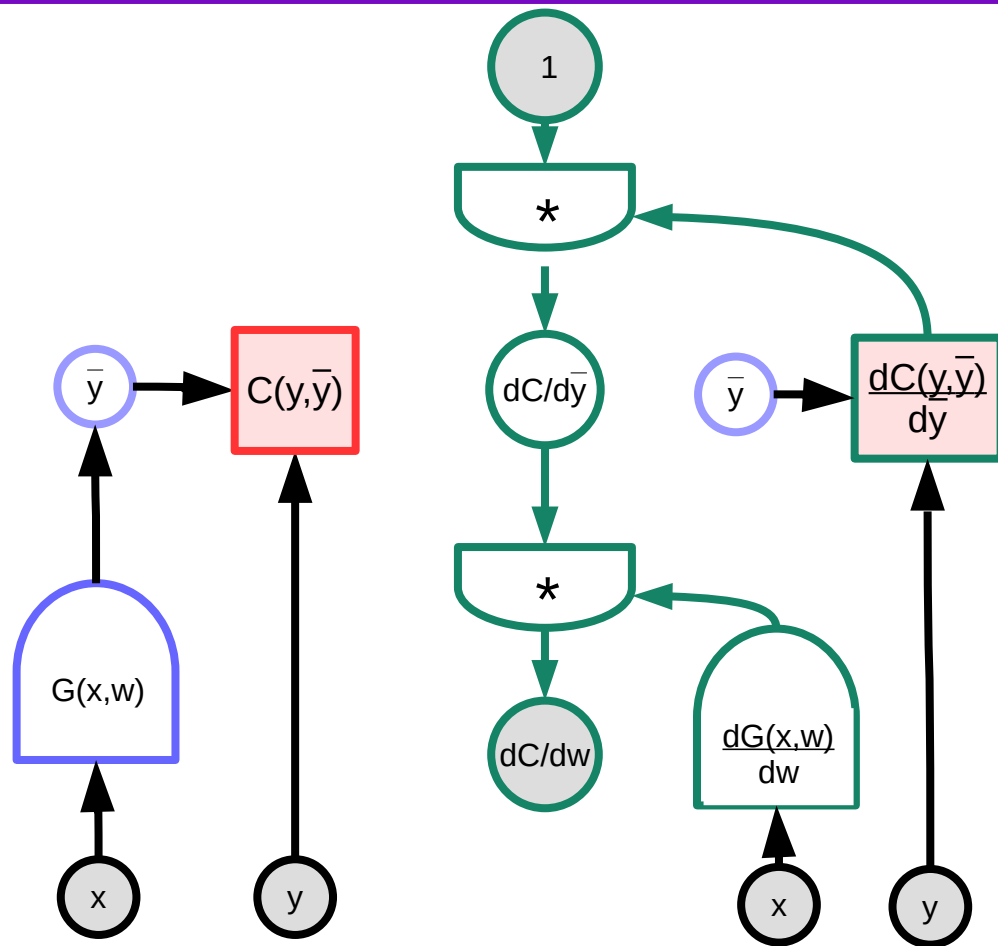
$$[1 \times N] = [1 \times M] \cdot [M \times N]$$

► Row vector = row vector . matrix

$$\frac{\partial C(y, \bar{y})}{\partial w} = \frac{\partial C(y, \bar{y})}{\partial \bar{y}} \frac{\partial G(x, w)}{\partial w}$$

$$[1 \times N] = [1 \times M] \cdot [M \times N]$$

► Gradient = gradient . Jacobian



Basic Modules

Linear

$$Y = W.X \quad ; \quad dC/dX = W^T \cdot dC/dY \quad ; \quad dC/dW = dC/dY \cdot X^T$$

ReLU

$$y = \text{ReLU}(x) \quad ; \quad \text{if } (x < 0) \quad dC/dx = 0 \quad \text{else} \quad dC/dx = dC/dy$$

Duplicate

$$Y1 = X, Y2 = X \quad ; \quad dC/dX = dC/dY1 + dC/dY2$$

Add

$$Y = X1 + X2 \quad ; \quad dC/dX1 = dC/dY \quad ; \quad dC/dX2 = dC/dY$$

Max

$$y = \max(x1, x2) \quad ; \quad \text{if } (x1 > x2) \quad dC/dx1 = dC/dy \quad \text{else} \quad dC/dx1 = 0$$

LogSoftMax

$$Y_i = X_i - \log[\sum_j \exp(X_j)] \quad ; \quad \dots???$$

Backprop in Practice

- Use ReLU non-linearities (tanh and logistic are falling out of favor)
- Use cross-entropy loss for classification
- Use Stochastic Gradient Descent on minibatches
- Shuffle the training samples
- Normalize the input variables (zero mean, unit variance)
- Schedule to decrease the learning rate
- Use a bit of L1 or L2 regularization on the weights (or a combination)
 - ▶ But it's best to turn it on after a couple of epochs
- Use “dropout” for regularization
 - ▶ Hinton et al 2012 <http://arxiv.org/abs/1207.0580>
- Lots more in [LeCun et al. “Efficient Backprop” 1998]
- Lots, lots more in “Neural Networks, Tricks of the Trade” (2012 edition) edited by G. Montavon, G. B. Orr, and K-R Müller (Springer)

- ▶ As long as there is a partial order on the modules

