



NEW YORK UNIVERSITY

Energy-Based Models (part 1)

<http://bit.ly/DLSP20>

Yann LeCun


NYU - Courant Institute & Center for Data Science

Facebook AI Research

<http://yann.lecun.com>



TAs: Alfredo Canziani, Mark Goldstein

Deep Learning, NYU, Spring 2020




Energy-Based Models

- ▶ **Feed-forward nets use a finite number of steps to produce a single output.**
- ▶ **What if...**
 - ▶ The problem requires a complex computation to produce its output? (complex inference)
 - ▶ There are multiple possible outputs for a single input? (e.g. predicting future video frames)
- ▶ **Inference through constraint satisfaction**
 - ▶ Finding an output that satisfies constraints: e.g a linguistically correct translation or speech transcription.
 - ▶ Maximum likelihood inference in graphical models




Energy-Based Models (EBM)

- ▶ **Energy function $F(x,y)$ scalar-valued.**
 - ▶ Takes low values when y is compatible with x and higher values when y is less compatible with x
- ▶ **Inference:** find values of y that make $F(x,y)$ small.
 - ▶ There may be multiple solutions $\tilde{y} = \operatorname{argmin}_y F(x, y)$
- ▶ **Note:** the energy is used for **inference**, not for learning





- ▶ **Example**
 - ▶ Blue dots are data points




Energy-Based Model: implicit function

- ▶ A feed-forward model is an **explicit function** that computes y from x .
- ▶ An EBM is an **implicit function** that captures the dependency between x and y
- ▶ Multiple y can be compatible with a single x




Energy-Based Model: implicit function

- ▶ **Energy function that captures the x,y dependencies:**
 - ▶ Low energy near the data points. Higher energy everywhere else.
 - ▶ If y is continuous, F should be smooth and differentiable, so we can use gradient-based inference algorithms.





$$\check{y} = \operatorname{argmin}_y F(x, y)$$



Energy-Based Model: gradient-based inference


- ▶ If y is continuous
- ▶ We can use a gradient-based method for inference.

$$\tilde{y} = \operatorname{argmin}_y F(x, y)$$

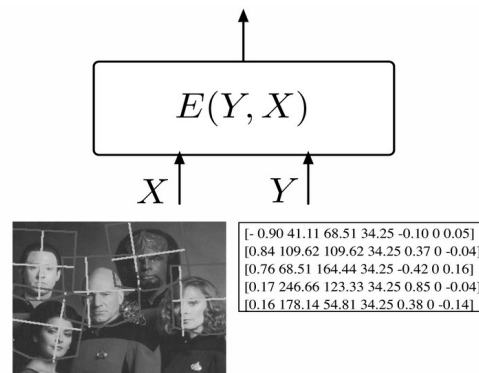


When inference is hard


- **Cases where inference is hard:**
 - Output is a high-dimensional object with structure:
 - Sequence, image, video,...
 - Output has compositional structure:
 - Text, action sequence,...
 - Output results from a long chain of reasoning
 - That can be reduced to an optimization problem




(a)




(b)




(c)



(d)



(e)




(f)

When inference involves latent variables

- ▶ **Latent variables are variables whose value is never given to us.**
- ▶ Examples: to read a handwritten word, it helps to know where the characters are



- ▶ To recognize speech, it helps to know where the words and phonemes are
 - ▶ You can read this if you understand english
 - ▶ Vous ne pouvez pas lire ceci si vous ne parlez pas français




When inference involves latent variables

- ▶ **Latent variables are variables whose value is never given to us.**
- ▶ Examples: to read a handwritten word, it helps to know where the characters are



- ▶ To recognize speech, it helps to know where the words and phonemes are
 - ▶ You can read this if you understand english
 - ▶ Vous ne pouvez pas lire ceci si vous ne parlez pas français



Latent-Variable EBM: inference

- ▶ Simultaneous minimization with respect to y and z


$$\check{y}, \check{z} = \operatorname{argmin}_{y,z} E(x, y, z)$$

- ▶ Redefinition of $F(x,y)$

$$F_{\infty}(x, y) = \operatorname{argmin}_z E(x, y, z)$$

$$F_{\beta}(x, y) = -\frac{1}{\beta} \log \int_z e^{-\beta E(x,y,z)}$$

$$\check{y} = \operatorname{argmin}_y F(x, y)$$




Latent-Variable EBM

- ▶ Allowing multiple predictions through a latent variable
- ▶ As z varies over a set, y varies over the manifold of possible predictions

$$F(x, y) = \min_z E(x, y, z)$$


- ▶ Useful then there are multiple correct (or plausible) outputs.
- ▶ Example: video prediction, text generation, translation, image synthesis....



Energy-Based Models vs Probabilistic Models

- ▶ **Probabilistic models are a special case of EBM**
 - ▶ Energies are like unnormalized negative log probabilities
- ▶ **Why use EBM instead of probabilistic models?**
 - ▶ EBM gives more flexibility in the choice of the scoring function.
 - ▶ More flexibility in the choice of objective function for learning
- ▶ **From energy to probability: Gibbs-Boltzmann distribution**
 - ▶ Beta is a positive constant

$$P(y|x) = \frac{1}{\beta} \frac{e^{-\beta F(x,y)}}{\int_{y'} e^{-\beta F(x,y')}}$$



Marginalizing over the latent variable


$$P(y, z|x) = \frac{e^{-\beta E(x,y,z)}}{\int_y \int_z e^{-\beta E(x,y,z)}} \quad P(y|x) = \int_z P(y, z|x)$$

$$P(y|x) = \frac{\int_z e^{-\beta E(x,y,z)}}{\int_y \int_z e^{-\beta E(x,y,z)}} = \frac{e^{-\beta \left[-\frac{1}{\beta} \log \int_z e^{-\beta E(x,y,z)} \right]}}{\int_y e^{-\beta \left[-\frac{1}{\beta} \log \int_z e^{-\beta E(x,y,z)} \right]}} = \frac{e^{-\beta F_\beta(x,y)}}{\int_y e^{\beta F_\beta(x,y)}}$$

► **Free energy $F(x,y)$** $F_\beta(x, y) = -\frac{1}{\beta} \log \int_z e^{-\beta E(x,y,z)}$

Self-Supervised Learning

Predict everything
from everything else



Self-Supervised Learning = Filling in the Blanks


- ▶ Predict any part of the input from any other part.

- ▶ Predict the **future** from the **past**.

- ▶ Predict the **masked** from the **visible**.

- ▶ Predict the **any occluded part** from **all available parts**.


- ▶ Pretend there is a part of the input you don't know and predict that.
- ▶ Reconstruction = SSL when any part could be known or unknown



Self-Supervised Learning: filling in the bl_nks


► Natural Language Processing: works great!

OUTPUT: This is a piece of text extracted from a large set of news articles



► Image Recognition / Understanding: works so-so

[Pathak et al 2014]



Learning Representations through Pretext SSL Tasks

- ▶ **Text / symbol sequences (discrete, works great!)**

- ▶ Future word(s) prediction (NLM)
- ▶ Masked words prediction (BERT et al.)

- ▶ **Image (continuous)**

- ▶ Inpainting, colorization, super-resolution

- ▶ **Video (continuous)**

- ▶ Future frame(s) prediction
- ▶ Masked frames prediction

- ▶ **Signal / Audio (continuous)**

- ▶ Restoration
- ▶ Future prediction

Self-Supervised Learning works **very** well for text

► Word2vec

► [Mikolov 2013]

► FastText

► [Joulin 2016] (FAIR)

► BERT

► Bidirectional Encoder Representations from Transformers

► [Devlin 2018]

► Cloze-Driven Auto-Encoder


► [Baevski 2019] (FAIR)

► RoBERTa [Ott 2019] (FAIR)

Use the output of the masked word's position to predict the masked word

Possible classes:
All English words


0.1%	Aardvark
...	...
10%	Improvisation
...	...
0%	Zyzyva



Randomly mask 15% of tokens

Input

Figure credit: Jay Alammar <http://jalammar.github.io/illustrated-bert/>



SSL works less well for images and video



input



Barnes et al. | 2009



Darabi et al. | 2012



Huang et al. | 2014



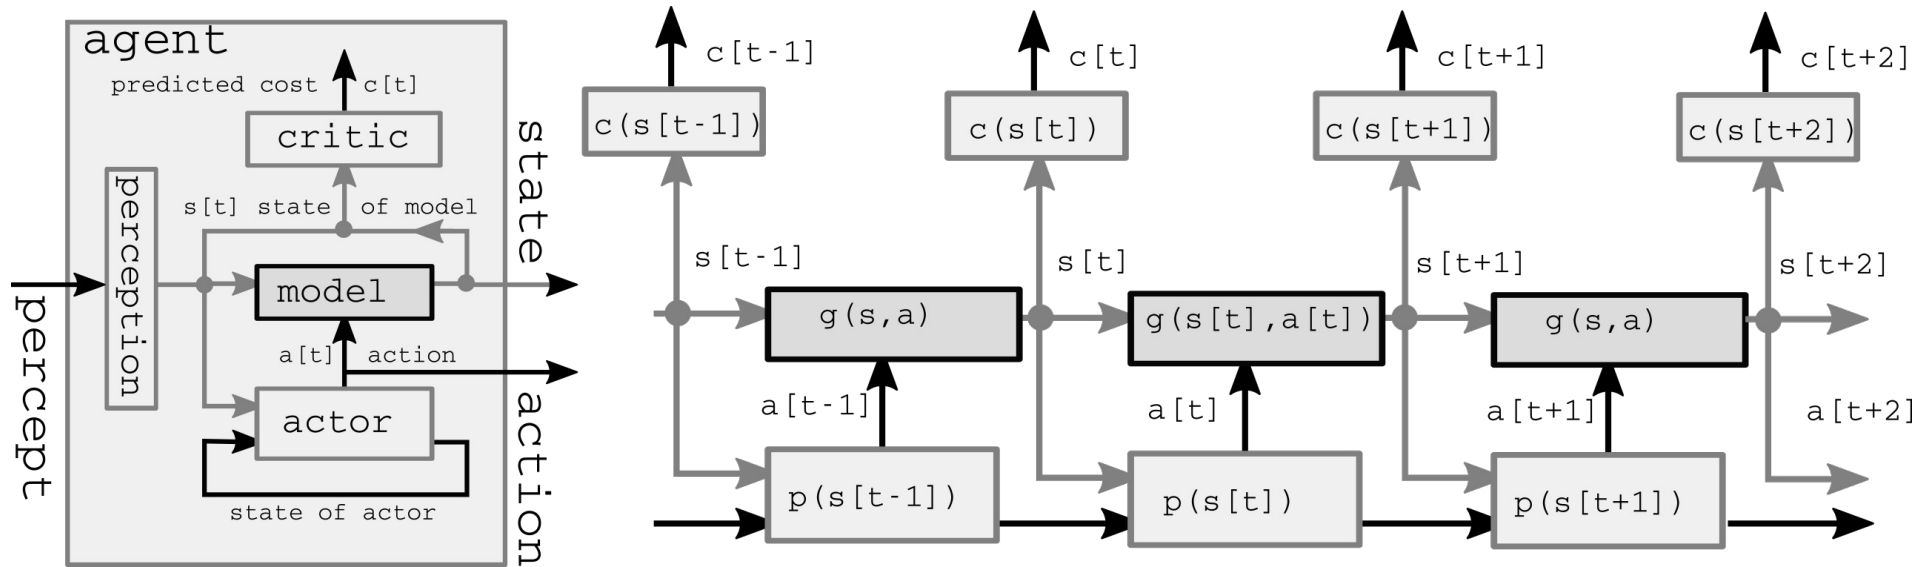
Pathak et al. | 2016



Iizuka et al. | 2017

Learning World Models for Autonomous AI Agents

- ▶ Learning **forward models** for control
- ▶ $s[t+1] = g(s[t], a[t], z[t])$
- ▶ Model-predictive control, model-predictive policy learning, model-based RL
- ▶ Robotics, games, dialog, HCI, etc




Self-Supervised Learning for Video Prediction

- ▶ The world is not entirely predictable
- ▶ There are many plausible continuations to a video segment



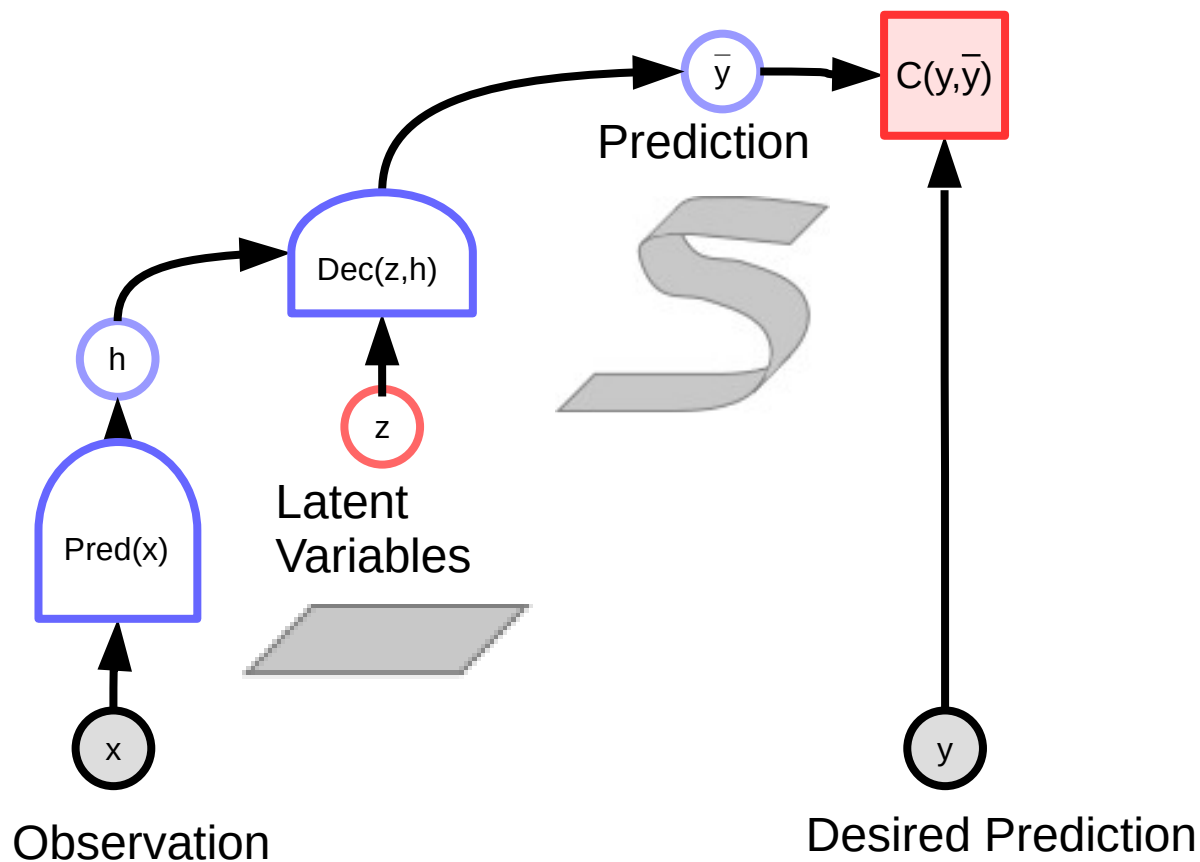
The world is stochastic

- ▶ Training a system to make a single prediction makes it predict the average of all plausible predictions
- ▶ Blurry predictions!



Solution: latent variable energy-based models

- ▶ Latent variables allows system to make multiple predictions



Self-supervised Adversarial Learning for Video Prediction


- ▶ Our brains are “prediction machines”
- ▶ Can we train machines to predict the future?
- ▶ Some success with “adversarial training”
 - ▶ [Mathieu, Couprie, LeCun arXiv:1511:05440]
- ▶ But we are far from a complete solution.



Problem: uncertainty!


- ▶ There are **many** plausible words that complete a text.
- ▶ There are **infinitely many** plausible frames to complete a video.
- ▶ **Deterministic predictors don't work!**
- ▶ **How to deal with uncertainty in the prediction?**

$$E(x, y) = C(y, G(x))$$



The world is not entirely predictable / stochastic

- ▶ **Video prediction:**
 - ▶ A deterministic predictor with L2 distance will predict the average of all plausible futures.
- ▶ **Blurry prediction!**




Energy-Based Model


► Scalar-valued energy function: $F(x,y)$

- measures the compatibility between x and y
- Low energy: y is good prediction from x
- High energy: y is bad prediction from x

- Inference: $\check{y} = \operatorname{argmin}_y F(x, y)$




Dark = low energy (good)
Bright = high energy (bad)
Purple = data manifold




[Figure from M-A Ranzato's PhD thesis]

Energy-Based Model: unconditional version

- ▶ **Scalar-valued energy function: $F(y)$**
 - ▶ measures the compatibility between the components of y
 - ▶ If we don't know in advance which part of y is known and which part is unknown
 - ▶ Example: auto-encoders, generative models (energy = $-\log$ likelihood)





Dark = low energy (good)
Bright = high energy (bad)
Purple = data manifold



Training an Energy-Based Model

- ▶ Parameterize $F(x,y)$
- ▶ Get training data $(x[i], y[i])$
- ▶ Shape $F(x,y)$ so that:
 - ▶ $F(x[i], y[i])$ is strictly smaller than $F(x[i], y)$ for all y different from $y[i]$
 - ▶ F is smooth (probabilistic methods break that!)
- ▶ **Two classes of learning methods:**
 - ▶ 1. **Contrastive methods:** push down on $F(x[i], y[i])$, push up on other points $F(x[i], y')$
 - ▶ 2. **Architectural Methods:** build $F(x,y)$ so that the volume of low energy regions is limited or minimized through regularization



Seven Strategies to Shape the Energy Function

- ▶ **Contrastive:** [they all are different ways to pick which points to push up]
 - ▶ C1: push down of the energy of data points, push up everywhere else: Max likelihood (needs tractable partition function or variational approximation)
 - ▶ C2: push down of the energy of data points, push up on chosen locations: max likelihood with MC/MMC/HMC, Contrastive divergence, Metric learning, Ratio Matching, Noise Contrastive Estimation, Min Probability Flow, adversarial generator/GANs
 - ▶ C3: train a function that maps points off the data manifold to points on the data manifold: denoising auto-encoder, masked auto-encoder (e.g. BERT)
- ▶ **Architectural:** [they all are different ways to limit the information capacity of the code]
 - ▶ A1: build the machine so that the volume of low energy stuff is bounded: PCA, K-means, Gaussian Mixture Model, Square ICA...
 - ▶ A2: use a regularization term that measures the volume of space that has low energy: Sparse coding, sparse auto-encoder, LISTA, Variational auto-encoders
 - ▶ A3: $F(x,y) = C(y, G(x,y))$, make $G(x,y)$ as "constant" as possible with respect to y : Contracting auto-encoder, saturating auto-encoder
 - ▶ A4: minimize the gradient and maximize the curvature around data points: score matching


Simple examples: PCA and K-means

■ Limit the capacity of z so that the volume of low energy stuff is bounded

► PCA, K-means, GMM, square ICA...

PCA: z is low dimensional


$$F(Y) = \|W^T W Y - Y\|^2$$



K-Means,

Z constrained to 1-of- K code

$$F(Y) = \min_z \sum_i \|Y - W_i Z_i\|^2$$



Familiar Example: Maximum Likelihood Learning

■ The energy can be interpreted as an unnormalized negative log density

■ Gibbs distribution: Probability proportional to $\exp(-\text{energy})$


▶ Beta parameter is akin to an inverse temperature

■ Don't compute probabilities unless you absolutely have to

▶ Because the denominator is often intractable

$$P(y) = \frac{\exp[-\beta F(y)]}{\int_{y'} \exp[-\beta F(y')]$$

$$P(y|x) = \frac{\exp[-\beta F(x, y)]}{\int_{y'} \exp[-\beta F(x, y')]$$



push down of the energy of data points, push up everywhere else


Max likelihood (requires a tractable partition function)

Maximizing $P(Y|W)$ on training samples

$$P(Y|W) = \frac{e^{-\beta E(Y,W)}}{\int_y e^{-\beta E(y,W)}} \quad \begin{array}{l} \text{make this big} \\ \text{make this small} \end{array}$$

Minimizing $-\log P(Y,W)$ on training samples

$$L(Y, W) = E(Y, W) + \frac{1}{\beta} \log \int_y e^{-\beta E(y, W)} \quad \begin{array}{l} \text{make this small} \\ \text{make this big} \end{array}$$



push down of the energy of data points, push up everywhere else

Gradient of the negative log-likelihood loss for one sample Y :


$$\frac{\partial L(Y, W)}{\partial W} = \frac{\partial E(Y, W)}{\partial W} - \int_y P(y|W) \frac{\partial E(y, W)}{\partial W}$$

Gradient descent:

$$W \leftarrow W - \eta \frac{\partial L(Y, W)}{\partial W}$$

Pushes down on the energy of the samples

Pulls up on the energy of low-energy Y 's



$$W \leftarrow W - \eta \frac{\partial E(Y, W)}{\partial W} + \eta \int_y P(y|W) \frac{\partial E(y, W)}{\partial W}$$


Latent-Variable EBM

- ▶ Allowing multiple predictions through a latent variable

- ▶ **Conditional:**

$$F(x, y) = \min_z E(x, y, z)$$


$$F(x, y) = -\frac{1}{\beta} \log \left[\int_z \exp(-\beta E(x, y, z)) \right]$$



- ▶ **Unconditional**

$$F(y) = \min_z E(y, z)$$

$$F(y) = -\frac{1}{\beta} \log \left[\int_z \exp(-\beta E(y, z)) \right]$$



Latent-Variable EBM for multimodal prediction


- ▶ Allowing multiple predictions through a latent variable
- ▶ As z varies over a set, y varies over the manifold of possible predictions

$$F(x, y) = \min_z E(x, y, z)$$

- ▶ **Examples:**

- ▶ K-means
- ▶ Sparse modeling
- ▶ GLO

[Bojanowski arXiv:1707.05776]





Latent-Variable EBM example: K-means

- ▶ Decoder is linear, z is a 1-hot vector (discrete)
- ▶ Energy function: $E(y, z) = \|y - Wz\|^2 \quad z \in 1 \text{ hot}$
- ▶ Inference by exhaustive search


$$F(y) = \min_z E(y, z)$$

- ▶ Volume of low-energy regions limited by number of prototypes k



Contrastive Embedding

- ▶ Distance measured in feature space
- ▶ Multiple “predictions” through feature invariance
- ▶ Siamese nets, metric learning [YLC NIPS'93,CVPR'05,CVPR'06]
- ▶ **Advantage: no pixel-level reconstruction**
- ▶ **Difficulty: hard negative mining**
- ▶ **Successful examples for images:**
 - ▶ DeepFace [Taigman et al. CVPR'14]
 - ▶ PIRL [Misra et al. To appear]
 - ▶ MoCo [He et al. Arxiv:1911.05722]
- ▶ **Video / Audio**
 - ▶ Temporal proximity [Taylor CVPR'11]
 - ▶ Slow feature [Goroshin NIPS'15]




Positive pair:
Make F small



Negative pair:
Make F large



MoCo on ImageNet [He et al. Arxiv:1911.05722]



Denoising AE: discrete

► [Vincent et al. JMLR 2008]

► **Masked Auto-Encoder**

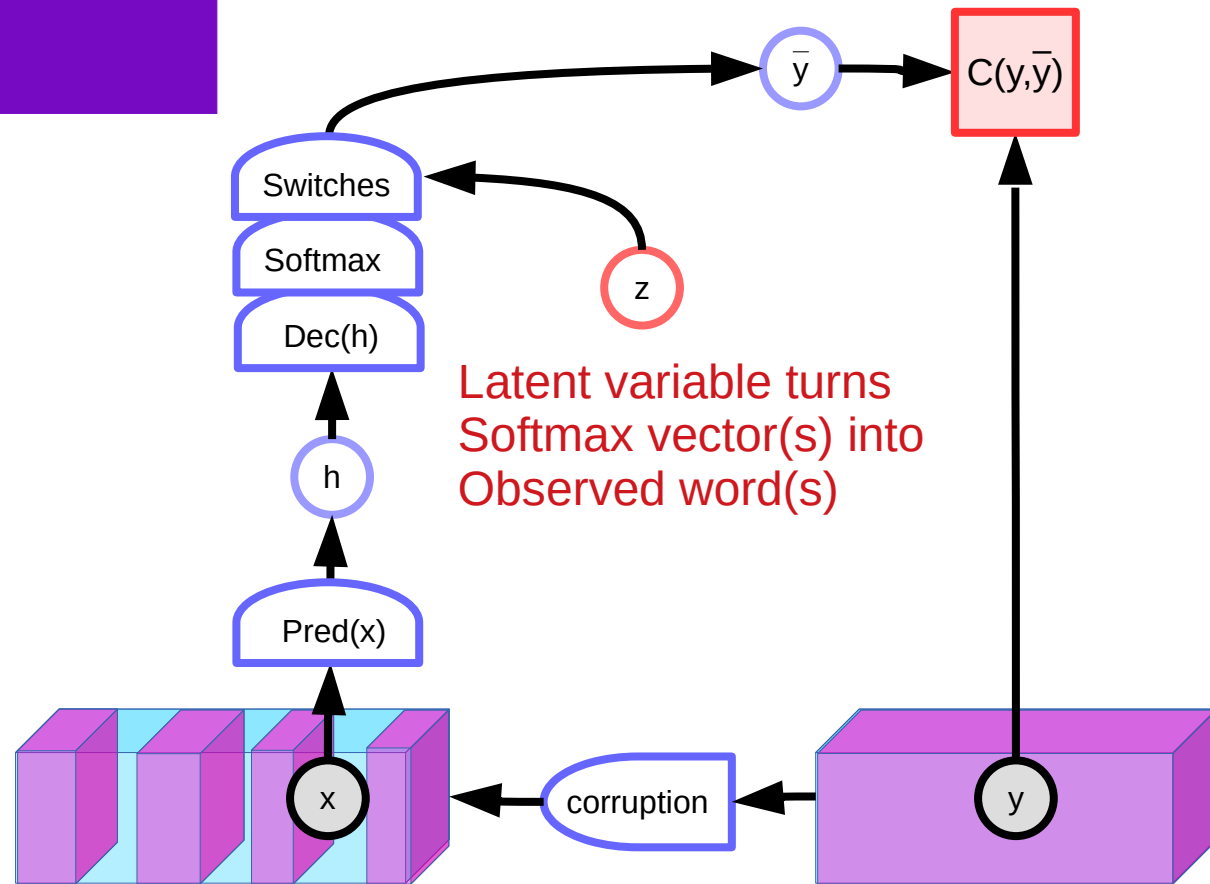
► [BERT et al.]

► **Issues:**

► latent variables are in output space

► No abstract LV to control the output

► How to cover the space of corruptions?

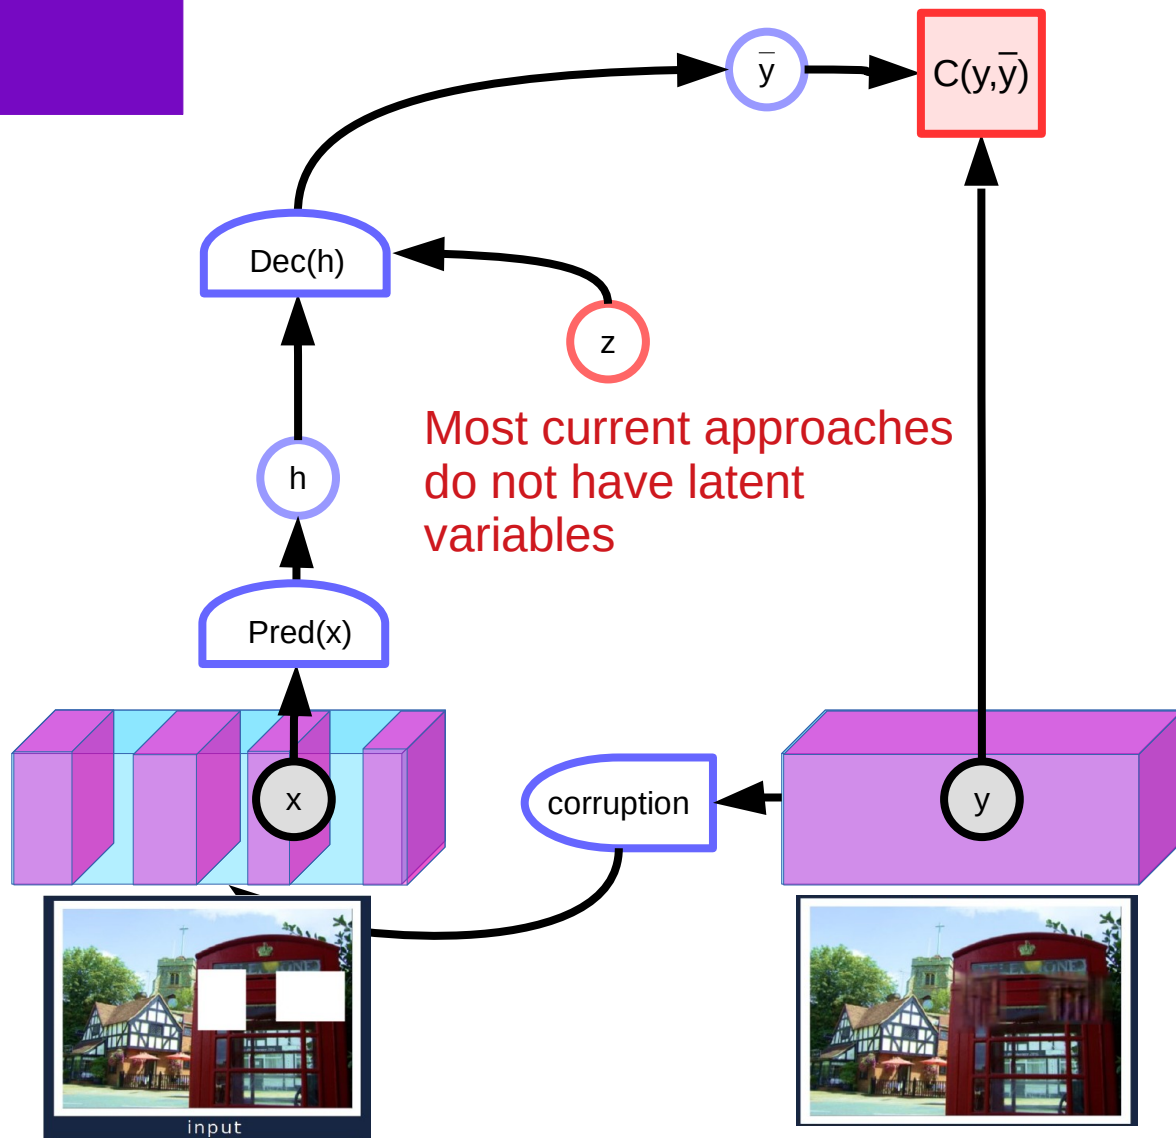



This is a [...] of text extracted
[...] a large set of [...] articles

This is a piece of text extracted
from a large set of news articles

Denoising AE: continuous

- Image inpainting [Pathak 17]
- Latent variables? GAN?




Prediction with Latent Variables

- ▶ If the Latent has too much capacity...
 - ▶ e.g. if it has the same dimension as y
- ▶ ... then the entire y space could be perfectly reconstructed

$$E(x, y, z) = C(y, \text{Dec}(\text{Pred}(x), z))$$

- ▶ For every y , there is always a z that will reconstruct it perfectly
 - ▶ The energy function would be zero everywhere
 - ▶ This is not a good model....
- ▶ **Solution: limiting the information capacity of the latent variable z .**



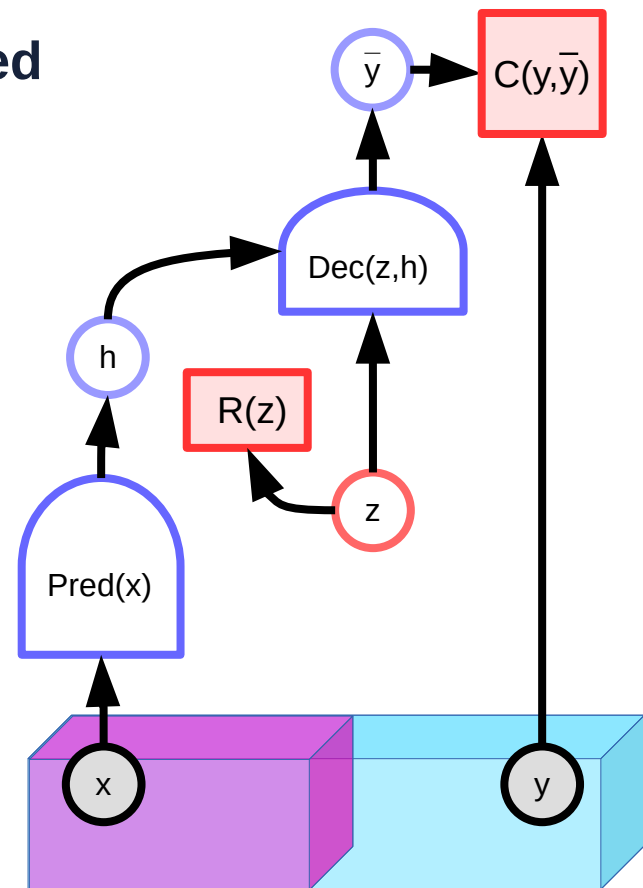
Regularized Latent Variable EBM

- ▶ Regularizer $R(z)$ limits the information capacity of z
- ▶ Without regularization, every y may be reconstructed exactly (flat energy surface)

$$E(x, y, z) = C(y, \text{Dec}(\text{Pred}(x), z)) + \lambda R(z)$$


- ▶ **Examples of $R(z)$:**

- ▶ Effective dimension
- ▶ Quantization / discretization
- ▶ L0 norm (# of non-0 components)
- ▶ L1 norm with decoder normalization
- ▶ Maximize lateral inhibition / competition
- ▶ Add noise to z while limiting its L2 norm (VAE)
- ▶ <your_information_throttling_method_goes_here>



Sequence → Abstract Features

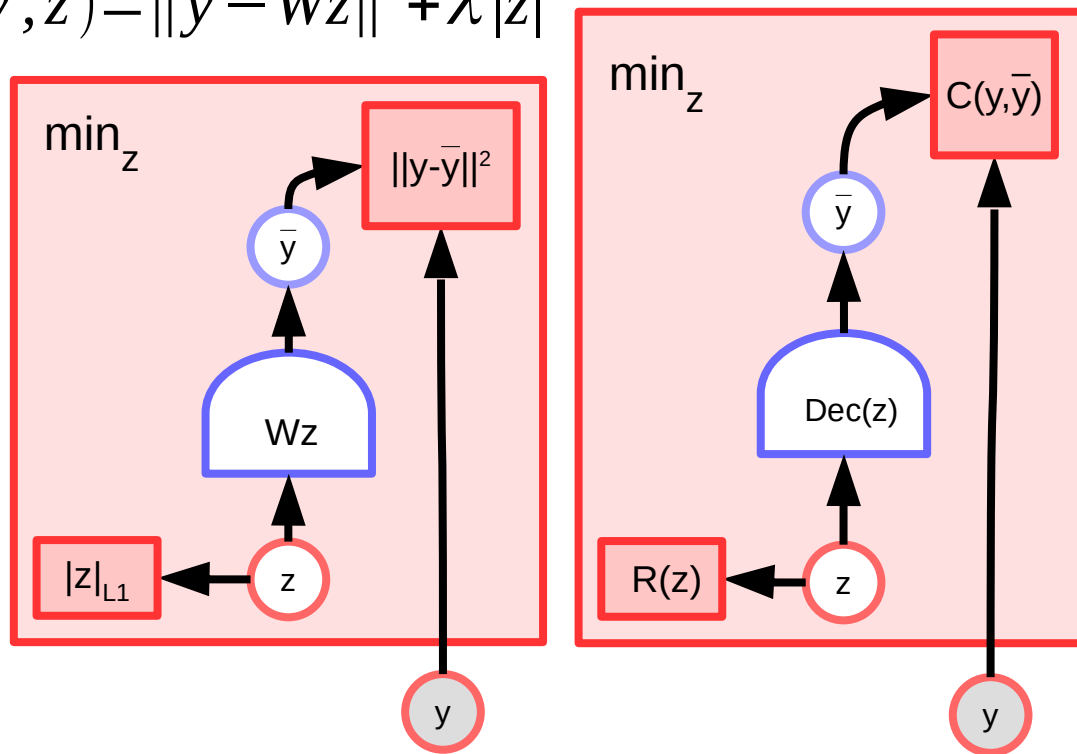

- ▶ Regularized LV EBM is passed over a sequence (e.g. a video, audio, text)
- ▶ The sequence of corresponding h and z is collected
- ▶ It contains all the information about the input sequence
- ▶ h contains the information in x that is useful to predict y
- ▶ z contains the complementary information, not present in x or h .
- ▶ Several such SSL modules can be stacked to learn hierarchical representations of sequences



Unconditional Regularized Latent Variable EBM

- Unconditional form. Reconstruction. No x , no predictor.
- Example: sparse modeling
 - Linear decoder
 - L1 regularizer on z

$$E(y, z) = \|y - Wz\|^2 + \lambda |z|$$



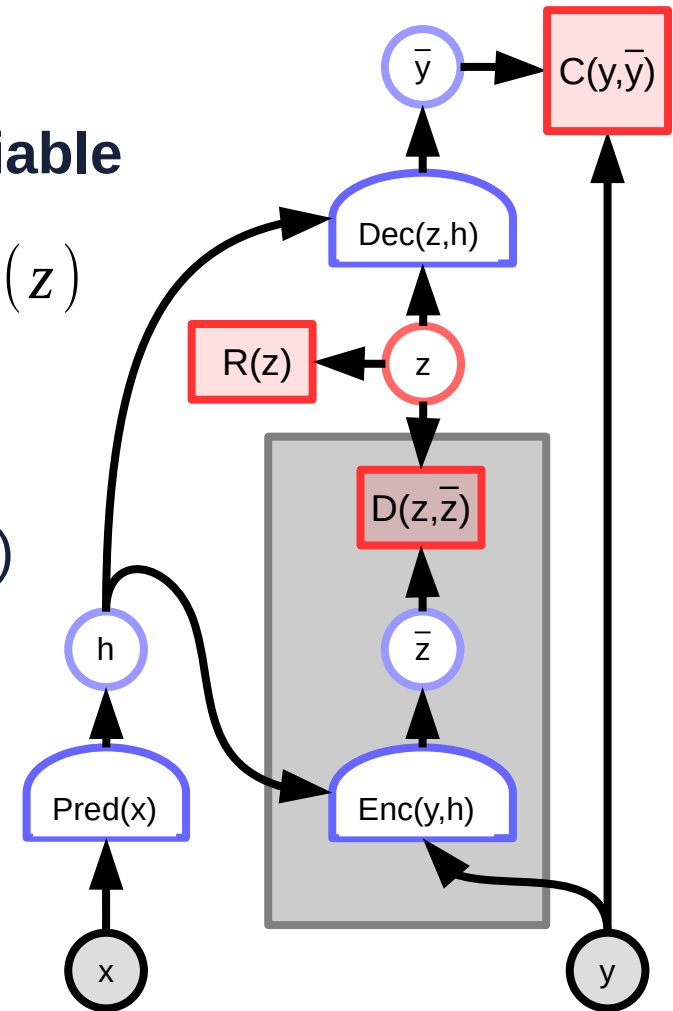

LatVar inference is expensive!

- ▶ Let's train an encoder to predict the latent variable

$$E(x, y, z) = C(y, \text{Dec}(z, h)) + D(z, \text{Enc}(x, y)) + \lambda R(z)$$


- ▶ Predictive Sparse Modeling

- ▶ $R(z)$ = L1 norm of z
- ▶ $\text{Dec}(z, h)$ gain must be bounded (clipped weights)
- ▶ Sparse Auto-Encoder
- ▶ LISTA [Gregor ICML 2010]




Sparse AE on handwritten digits (MNIST)

- ▶ **256 basis functions** Basis functions (columns of decoder matrix) are digit parts
- ▶ **All digits are a linear combination of a small number of these**




Predictive Sparse Decomposition (PSD): Training

- ▶ **Training on natural images patches.**
- ▶ 12X12
- ▶ 256 basis functions
- ▶ [Ranzato 2007]



iteration no 0

Learned Features: V1-like receptive fields




Convolutional Sparse Auto-Encoder on Natural Images

- **Filters and Basis Functions obtained. Linear decoder (conv)**
 - with 1, 2, 4, 8, 16, 32, and 64 filters [Kavukcuoglu NIPS 2010]


Encoder Filters

Decoder Filters



Encoder Filters


Decoder Filters




Convolutional Sparse Auto-Encoder on Natural Images

- ▶ Trained on CIFAR 10 (32x32 color images)
- ▶ Architecture: Linear decoder, LISTA recurrent encoder
- ▶ Pytorch implementation (talk to Jure Zbontar)

sparse codes (z) from encoder

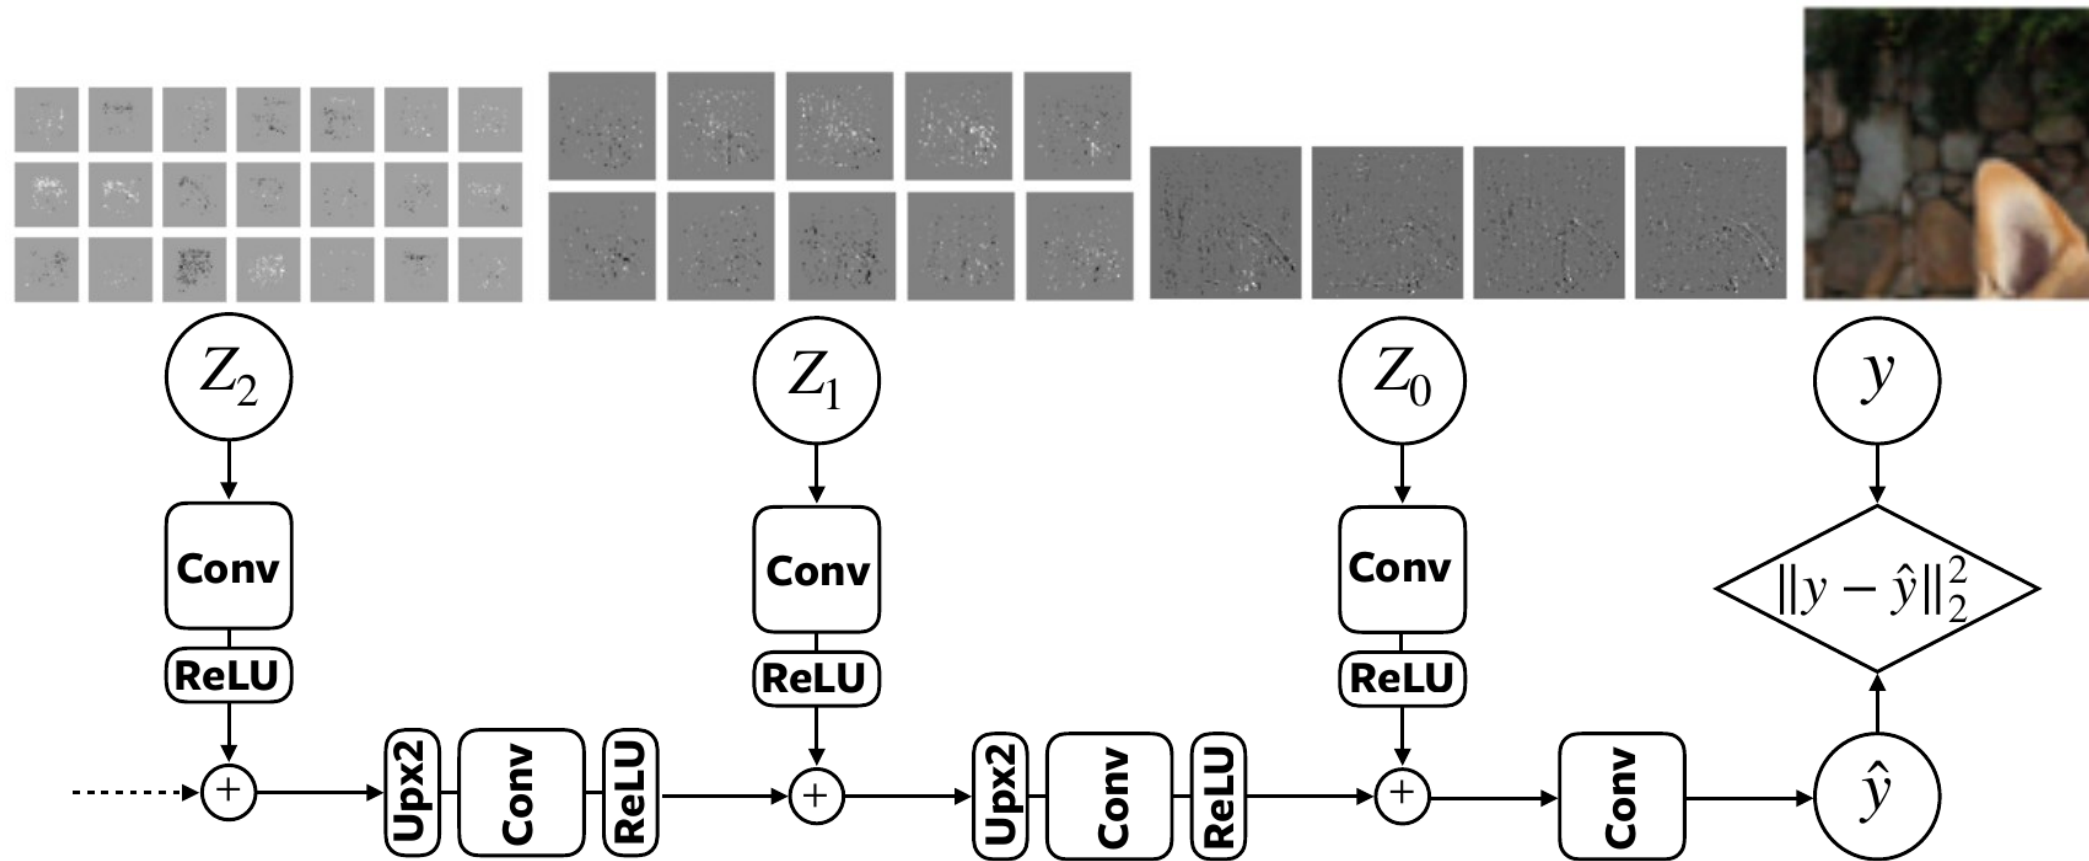


9x9 decoder kernels



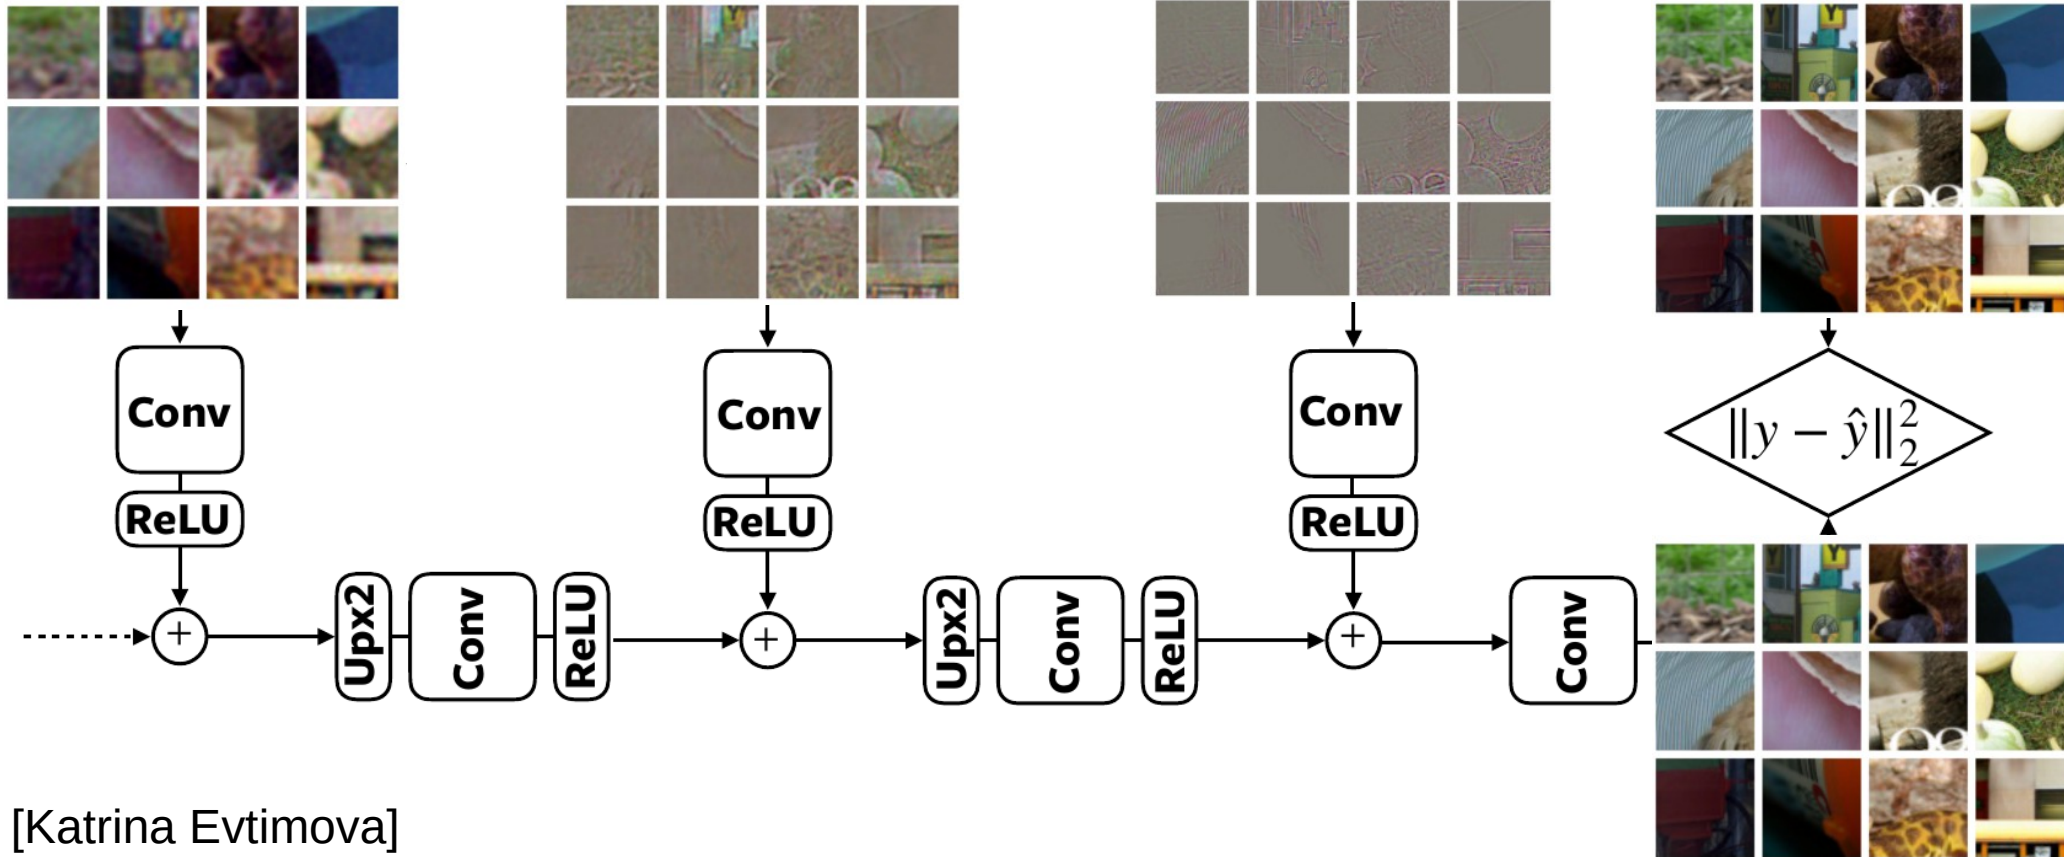
Multilayer Convolutional Sparse Modeling

► Learning hierarchical representations



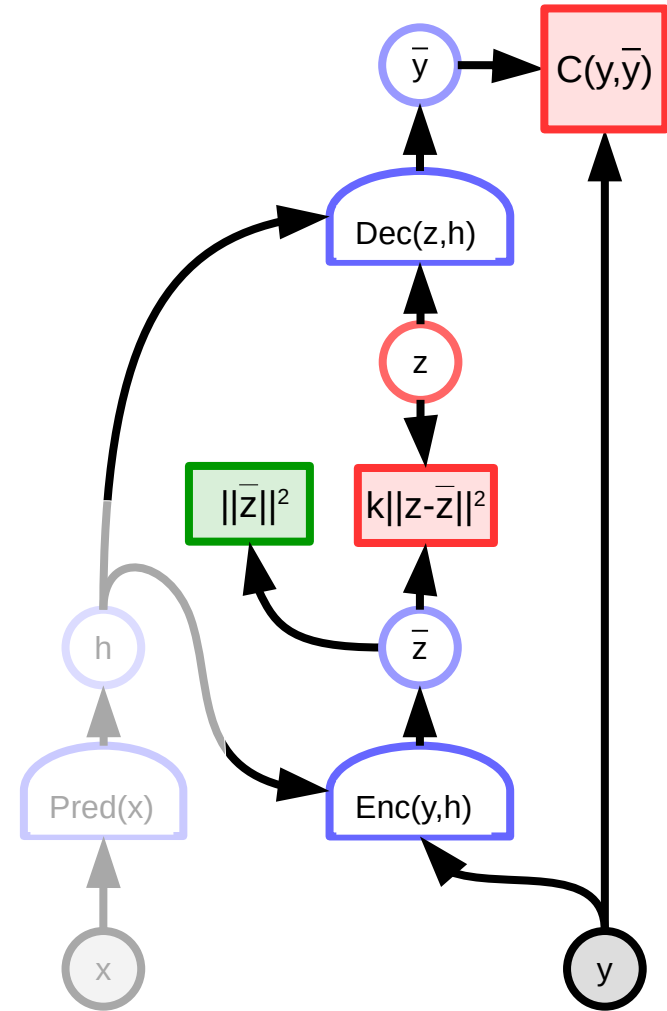
Multilayer Convolutional Sparse Modeling

► Reconstructions from Z2, Z1, Z0 and all of (Z2,Z1,Z0)




Variational Auto-Encoder

- ▶ Limiting the information capacity of the code by adding Gaussian noise
- ▶ The energy term $k||z-\bar{z}||^2$ is seen as the log of a prior from which to sample z
- ▶ The encoder output is regularized to have a mean and a variance close to zero.




Variational Auto-Encoder

► Code vectors for training samples




Variational Auto-Encoder

- ▶ **Code vectors for training sample with Gaussian noise**
 - ▶ Some fuzzy balls overlap, causing bad reconstructions



Variational Auto-Encoder

- ▶ The code vectors want to move away from each other to minimize reconstruction error
- ▶ But that does nothing for us



Variational Auto-Encoder

- ▶ **Attach the balls to the center with a spring, so they don't fly away**
 - ▶ Minimize the square distances of the balls to the origin
- ▶ **Center the balls around the origin**
 - ▶ Make the center of mass zero
- ▶ **Make the sizes of the balls close to 1 in each dimension**
 - ▶ Through a so-called KL term

